

Housing Prices in King County

Author: Carlos Garza

Overview

This notebook contains a regression analysis of the cost of King County real estate. Utilizing the CRISP-DM framework, linear regression models, and statistical techniques, I created and refined a model that describes the cost of real estate in King County in relation to a list of independent variables.

My data, methodology, and derived conclusions are detailed in the body of this document.

Business Problem

To gain an edge in the industry, a Seattle-based real estate company wants to automate their initial appraisal process. Developing an algorithm to accurately appraise the value of a house without physically inspecting the property can be an invaluable advantage in the fast paced real estate market of a rapidly expanding city. Using county data and multiple linear regression I created a model the company can use to predict real estate values.

Data

The data utilized in this model describes houses sold in 2014 and 2015.

The data is summarized below.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('darkgrid')
%matplotlib inline
```

```
In [2]: df = pd.read_csv('data/kc_house_data.csv')
df.head()
```

```
Out[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfr
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	N
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

```
In [3]: # Data Summary
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  float64
9   view                   21534 non-null  float64
10  condition              21597 non-null  int64
11  grade                  21597 non-null  int64
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                   21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

```
In [4]: # Exploration of Values
for column in df.columns:
    print(column, '\n')
    print(df[column].value_counts())
    print('_____')
```

id

```
795000620      3
1825069031      2
2019200220      2
7129304540      2
1781500435      2
..
7812801125      1
4364700875      1
3021059276      1
880000205       1
1777500160      1
Name: id, Length: 21420, dtype: int64
```

date

```
6/23/2014      142
6/25/2014      131
6/26/2014      131
7/8/2014       127
4/27/2015      126
...
2/15/2015       1
8/3/2014        1
1/17/2015       1
1/10/2015       1
```

```
5/24/2015      1
Name: date, Length: 372, dtype: int64
```

price

```
350000.0      172
450000.0      172
550000.0      159
500000.0      152
425000.0      150
...
870515.0       1
336950.0       1
386100.0       1
176250.0       1
884744.0       1
Name: price, Length: 3622, dtype: int64
```

bedrooms

```
3      9824
4      6882
2      2760
5      1601
6       272
1       196
7        38
8        13
9         6
10        3
11         1
33         1
Name: bedrooms, dtype: int64
```

bathrooms

```
2.50      5377
1.00      3851
1.75      3048
2.25      2047
2.00      1930
1.50      1445
2.75      1185
3.00       753
3.50       731
3.25       589
3.75       155
4.00       136
4.50       100
4.25        79
0.75        71
4.75        23
5.00        21
5.25        13
5.50        10
1.25         9
6.00         6
5.75         4
0.50         4
8.00         2
6.25         2
6.75         2
6.50         2
7.50         1
7.75         1
```

Name: bathrooms, dtype: int64

sqft_living

1300	138
1400	135
1440	133
1660	129
1010	129

...

4970	1
2905	1
2793	1
4810	1
1975	1

Name: sqft_living, Length: 1034, dtype: int64

sqft_lot

5000	358
6000	290
4000	251
7200	220
7500	119

...

1448	1
38884	1
17313	1
35752	1
315374	1

Name: sqft_lot, Length: 9776, dtype: int64

floors

1.0	10673
2.0	8235
1.5	1910
3.0	611
2.5	161
3.5	7

Name: floors, dtype: int64

waterfront

0.0	19075
1.0	146

Name: waterfront, dtype: int64

view

0.0	19422
2.0	957
3.0	508
1.0	330
4.0	317

Name: view, dtype: int64

condition

3	14020
4	5677
5	1701
2	170
1	29

Name: condition, dtype: int64

grade

7	8974
8	6065
9	2615
6	2038
10	1134
11	399
5	242
12	89
4	27
13	13
3	1

Name: grade, dtype: int64

sqft_above

1300	212
1010	210
1200	206
1220	192
1140	184
...	
2601	1
440	1
2473	1
2441	1
1975	1

Name: sqft_above, Length: 942, dtype: int64

sqft_basement

0.0	12826
?	454
600.0	217
500.0	209
700.0	208
...	
2180.0	1
3480.0	1
2400.0	1
2050.0	1
1770.0	1

Name: sqft_basement, Length: 304, dtype: int64

yr_built

2014	559
2006	453
2005	450
2004	433
2003	420
...	
1933	30
1901	29
1902	27
1935	24
1934	21

Name: yr_built, Length: 116, dtype: int64

yr_renovated

0.0	17011
2014.0	73

2003.0	31
2013.0	31
2007.0	30

...

1946.0	1
1959.0	1
1971.0	1
1951.0	1
1954.0	1

Name: yr_renovated, Length: 70, dtype: int64

zipcode

98103	602
98038	589
98115	583
98052	574
98117	553

...

98102	104
98010	100
98024	80
98148	57
98039	50

Name: zipcode, Length: 70, dtype: int64

lat

47.6624	17
47.5491	17
47.5322	17
47.6846	17
47.6711	16

..

47.2785	1
47.4162	1
47.3870	1
47.2313	1
47.2715	1

Name: lat, Length: 5033, dtype: int64

long

-122.290	115
-122.300	111
-122.362	104
-122.291	100
-122.372	99

...

-121.403	1
-121.804	1
-121.726	1
-121.895	1
-121.893	1

Name: long, Length: 751, dtype: int64

sqft_living15

1540	197
1440	195
1560	192
1500	180
1460	169

...

4890	1
------	---

```

2873      1
952      1
3193      1
2049      1
Name: sqft_living15, Length: 777, dtype: int64

```

sqft_lot15

```

5000      427
4000      356
6000      288
7200      210
4800      145
...
11036      1
8989      1
871200      1
809      1
6147      1
Name: sqft_lot15, Length: 8682, dtype: int64

```

Data Preprocessing

To begin, columns we are not interested in will be dropped. We are not concerned with the date the house was sold or if the house has been viewed.

Additionally, because all of our data ranges over only two years, we will count multiple sales of the same house independently, and therefore can drop the house ID. The rational behind this is that houses constantly on the market may hint at some underlining issue that may or may not be described numerically by our data so we will leave each sale in our analysis to obtain the broadest picture.

```
In [5]: df.drop(columns = ['id', 'date', 'view'], inplace = True)
```

Observing the data types and value counts in our data exploration above it can be seen that basement square footage is stored as a string. most values can be expressed as floats, so we will change data types to be treated as a continuous variable. unknown values (stored as '?') will be changed to 0.0

```
In [6]: df['sqft_basement'].replace('?', '0.0', inplace = True)
df['sqft_basement'] = df['sqft_basement'].astype('float')
```

Now that all of our data is in a usable format, we can create our first baseline model

Baseline Model

To keep as much data as possible in all future models, we will not drop missing values from our data until just before creating a new model

```
In [7]: baseline_data = df.dropna()
baseline_data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 15809 entries, 1 to 21596
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype

```

```

-----
0  price          15809 non-null float64
1  bedrooms      15809 non-null int64
2  bathrooms     15809 non-null float64
3  sqft_living   15809 non-null int64
4  sqft_lot      15809 non-null int64
5  floors        15809 non-null float64
6  waterfront    15809 non-null float64
7  condition     15809 non-null int64
8  grade         15809 non-null int64
9  sqft_above    15809 non-null int64
10 sqft_basement 15809 non-null float64
11 yr_built      15809 non-null int64
12 yr_renovated  15809 non-null float64
13 zipcode       15809 non-null int64
14 lat           15809 non-null float64
15 long          15809 non-null float64
16 sqft_living15 15809 non-null int64
17 sqft_lot15    15809 non-null int64
dtypes: float64(8), int64(10)
memory usage: 2.3 MB

```

```
In [8]: import statsmodels.api as sm
```

```
In [9]: baseline_predictors = baseline_data.drop('price', axis=1)
model = sm.OLS(baseline_data['price'], sm.add_constant(baseline_predictors)).fit()
model.summary()
```

```
Out[9]:
```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.694
Model:	OLS	Adj. R-squared:	0.693
Method:	Least Squares	F-statistic:	2103.
Date:	Sat, 16 Jan 2021	Prob (F-statistic):	0.00
Time:	13:32:41	Log-Likelihood:	-2.1594e+05
No. Observations:	15809	AIC:	4.319e+05
Df Residuals:	15791	BIC:	4.320e+05
Df Model:	17		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-2.373e+06	3.51e+06	-0.675	0.499	-9.26e+06	4.51e+06
bedrooms	-4.066e+04	2249.400	-18.074	0.000	-4.51e+04	-3.62e+04
bathrooms	4.608e+04	3924.315	11.743	0.000	3.84e+04	5.38e+04
sqft_living	157.4150	21.879	7.195	0.000	114.530	200.300
sqft_lot	0.1354	0.057	2.371	0.018	0.023	0.247
floors	5905.1578	4336.868	1.362	0.173	-2595.600	1.44e+04
waterfront	7.917e+05	1.92e+04	41.132	0.000	7.54e+05	8.29e+05
condition	2.762e+04	2829.941	9.759	0.000	2.21e+04	3.32e+04
grade	9.804e+04	2603.022	37.663	0.000	9.29e+04	1.03e+05

	master_analysis					
sqft_above	35.3500	21.823	1.620	0.105	-7.425	78.125
sqft_basement	14.4558	21.658	0.667	0.504	-27.997	56.908
yr_built	-2821.6048	87.044	-32.416	0.000	-2992.221	-2650.988
yr_renovated	21.8909	4.438	4.932	0.000	13.191	30.590
zipcode	-496.6853	39.529	-12.565	0.000	-574.167	-419.204
lat	5.813e+05	1.28e+04	45.244	0.000	5.56e+05	6.07e+05
long	-2.313e+05	1.58e+04	-14.639	0.000	-2.62e+05	-2e+05
sqft_living15	27.9778	4.128	6.778	0.000	19.887	36.069
sqft_lot15	-0.3309	0.086	-3.847	0.000	-0.500	-0.162
Omnibus:	13670.728	Durbin-Watson:	1.972			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1430940.355			
Skew:	3.652	Prob(JB):	0.00			
Kurtosis:	49.032	Cond. No.	2.15e+08			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.15e+08. This might indicate that there are strong multicollinearity or other numerical problems.

Model Tuning/ Reiteration

Statistically Insignificant Variables/Collinearity

To achieve a better fit, the model can be tuned. Setting alpha to 0.05, we can determine which variables are not statistically significant and can be dropped. Additionally, examining the collinearity of the independent variables can give clues on what variables to drop to improve our R-squared value in the next iteration. The table below shows the correlation between the independent variables.

```
In [10]: lin_test=baseline_predictors.corr().abs().stack().reset_index().sort_values(0, a
# zip the variable name columns (which were only named level_0 and level_1 by de
lin_test['pairs'] = list(zip(lin_test.level_0, lin_test.level_1))
# set index to pairs
lin_test.set_index(['pairs'], inplace = True)
# drop level columns
lin_test.drop(columns=['level_1', 'level_0'], inplace = True)
# rename correlation column as cc rather than 0
lin_test.columns = ['cc']
# drop duplicates. This could be dangerous if you have variables perfectly corre
# for the sake of exercise, kept it in.
lin_test.drop_duplicates(inplace=True)
lin_test.head(15)
```

Out[10]:

cc

	pairs	cc
pairs		
(bedrooms, bedrooms)	1.000000	
(sqft_above, sqft_living)	0.876023	
(grade, sqft_living)	0.764699	
(grade, sqft_above)	0.758407	
(sqft_living, sqft_living15)	0.756818	
(sqft_living, bathrooms)	0.754361	
(sqft_above, sqft_living15)	0.732934	
(sqft_lot, sqft_lot15)	0.719935	
(sqft_living15, grade)	0.717371	
(sqft_above, bathrooms)	0.686171	
(grade, bathrooms)	0.665321	
(sqft_living, bedrooms)	0.573750	
(bathrooms, sqft_living15)	0.570180	
(zipcode, long)	0.562116	
(sqft_above, floors)	0.529702	

To start, variables with a P-value above our threshhold (alpha) are dropped. Additionally, some variables with high collinearity are dropped.

```
In [11]: to_drop = ['floors', 'sqft_above', 'sqft_basement', 'sqft_lot', 'sqft_lot15', 'sqft_living15']
df.drop(to_drop, axis=1, inplace=True)
df.head()
```

	price	bedrooms	bathrooms	sqft_living	waterfront	condition	grade	yr_built	yr_renovated
0	221900.0	3	1.00	1180	NaN	3	7	1955	1995
1	538000.0	3	2.25	2570	0.0	3	7	1951	1995
2	180000.0	2	1.00	770	0.0	3	6	1933	1995
3	604000.0	4	3.00	1960	0.0	5	7	1965	1995
4	510000.0	3	2.00	1680	0.0	3	8	1987	1995

```
In [12]: model_2_data = df.dropna()
model_2_predictors = model_2_data.drop('price', axis=1)
model = sm.OLS(model_2_data['price'], sm.add_constant(model_2_predictors)).fit()
model.summary()
```

OLS Regression Results			
Dep. Variable:	price	R-squared:	0.692
Model:	OLS	Adj. R-squared:	0.692

Method: Least Squares **F-statistic:** 3224.
Date: Sat, 16 Jan 2021 **Prob (F-statistic):** 0.00
Time: 13:32:41 **Log-Likelihood:** -2.1598e+05
No. Observations: 15809 **AIC:** 4.320e+05
Df Residuals: 15797 **BIC:** 4.321e+05
Df Model: 11
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	2.651e+06	3.36e+06	0.790	0.430	-3.93e+06	9.23e+06
bedrooms	-4.074e+04	2244.302	-18.154	0.000	-4.51e+04	-3.63e+04
bathrooms	4.565e+04	3783.151	12.068	0.000	3.82e+04	5.31e+04
sqft_living	194.4893	3.662	53.114	0.000	187.312	201.667
waterfront	7.928e+05	1.93e+04	41.112	0.000	7.55e+05	8.31e+05
condition	2.453e+04	2810.180	8.728	0.000	1.9e+04	3e+04
grade	1.071e+05	2407.087	44.478	0.000	1.02e+05	1.12e+05
yr_built	-2791.2713	84.465	-33.046	0.000	-2956.832	-2625.710
yr_renovated	20.4666	4.438	4.611	0.000	11.767	29.166
zipcode	-516.7809	39.280	-13.156	0.000	-593.774	-439.788
lat	5.81e+05	1.28e+04	45.511	0.000	5.56e+05	6.06e+05
long	-2.06e+05	1.5e+04	-13.746	0.000	-2.35e+05	-1.77e+05

Omnibus: 13471.346 **Durbin-Watson:** 1.969
Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 1339076.232
Skew: 3.580 **Prob(JB):** 0.00
Kurtosis: 47.515 **Cond. No.** 1.99e+08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.99e+08. This might indicate that there are strong multicollinearity or other numerical problems.

We will not drop any more full columns, so we can commit to dropping rows with missing data.

```
In [13]: df.dropna(inplace = True)
```

Categorical Variables

Next, we'll recheck the values of our columns to see if any variables are categorical.

```
In [14]: # Exploration of Values
for column in df.columns:
    print(column, '\n')
    print(df[column].value_counts())
    print('_____')
```

price

```
350000.0    130
450000.0    121
500000.0    115
550000.0    113
425000.0    111
...
959750.0     1
503500.0     1
927000.0     1
648475.0     1
311300.0     1
Name: price, Length: 3042, dtype: int64
```

bedrooms

```
3    7143
4    5094
2    2007
5    1186
6     194
1     141
7      23
8       10
9         6
10        3
11         1
33         1
Name: bedrooms, dtype: int64
```

bathrooms

```
2.50    4013
1.00    2769
1.75    2235
2.25    1494
2.00    1398
1.50    1064
2.75     853
3.50     544
3.00     544
3.25     432
3.75     104
4.00      100
4.50       75
4.25       62
0.75       50
4.75       17
5.00       14
5.25       11
5.50        8
6.00        6
1.25        6
0.50        3
8.00        2
5.75        2
7.75        1
```

```
7.50      1
6.75      1
Name: bathrooms, dtype: int64
```

sqft_living

```
1820      102
1440      100
1400       98
1300       95
1320       94
...
1794       1
1802       1
1834       1
5960       1
1715       1
Name: sqft_living, Length: 913, dtype: int64
```

waterfront

```
0.0      15688
1.0       121
Name: waterfront, dtype: int64
```

condition

```
3      10252
4       4151
5       1256
2        131
1         19
Name: condition, dtype: int64
```

grade

```
7      6562
8      4444
9      1927
6      1488
10     835
11     291
5      167
12     67
4       16
13      11
3         1
Name: grade, dtype: int64
```

yr_built

```
2014     401
2006     335
2005     328
2007     310
2004     307
...
1901      22
1933      18
1902      18
1934      15
1935      14
Name: yr_built, Length: 116, dtype: int64
```

yr_renovated

```

0.0      15157
2014.0     64
2013.0     29
2005.0     27
2000.0     25
...
1974.0      1
1959.0      1
1934.0      1
1944.0      1
1976.0      1
Name: yr_renovated, Length: 70, dtype: int64

```

```

zipcode
98038     439
98103     426
98052     416
98115     409
98042     409
...
98010      70
98102      65
98024      59
98148      42
98039      36
Name: zipcode, Length: 70, dtype: int64

```

```

lat
47.6955     14
47.6846     14
47.6647     13
47.6711     13
47.6624     13
..
47.7697      1
47.4195      1
47.3227      1
47.1903      1
47.2051      1
Name: lat, Length: 4754, dtype: int64

```

```

long
-122.290     86
-122.300     78
-122.365     78
-122.288     77
-122.357     75
..
-121.822      1
-121.739      1
-122.475      1
-121.833      1
-121.676      1
Name: long, Length: 729, dtype: int64

```

Condition, grade, and zipcode columns can likely be treated as categorical variables, so will be converted to dummy variables.

```
In [15]: categoricals = ['condition', 'grade', 'zipcode']
```

```

condition_cat = pd.get_dummies(df['condition'], prefix = 'condition', drop_first
grade_cat = pd.get_dummies(df['grade'], prefix = 'grade', drop_first=True)
zipcode_cat = pd.get_dummies(df['zipcode'], prefix = 'zipcode', drop_first = Tru
df = df.drop(columns = categoricals)
df = pd.concat([df, condition_cat, grade_cat, zipcode_cat], axis=1)

```

additionally, the 'yr_renovated' column will be converted to a categorical 'renovated' column

```

In [16]: df['renovated'] = df['yr_renovated'].map(lambda x: 0 if x == 0 else 1)
df.drop(columns=['yr_renovated'], inplace = True)
df.head()

```

```

Out[16]:

```

	price	bedrooms	bathrooms	sqft_living	waterfront	yr_built	lat	long	condition
1	538000.0	3	2.25	2570	0.0	1951	47.7210	-122.319	
3	604000.0	4	3.00	1960	0.0	1965	47.5208	-122.393	
4	510000.0	3	2.00	1680	0.0	1987	47.6168	-122.045	
5	1230000.0	4	4.50	5420	0.0	2001	47.6561	-122.005	
6	257500.0	3	2.25	1715	0.0	1995	47.3097	-122.327	

5 rows × 92 columns

Future models in this notebook will be created with Scikit Learn

```

In [17]: from sklearn.linear_model import LinearRegression

```

```

In [18]: lin_reg_3 = LinearRegression()
y = df['price']
X = df.drop(['price'], axis = 1)
lin_reg_3.fit(X, y)
print('R-Squared: ', lin_reg_3.score(X, y))

```

R-Squared: 0.8241407664576075

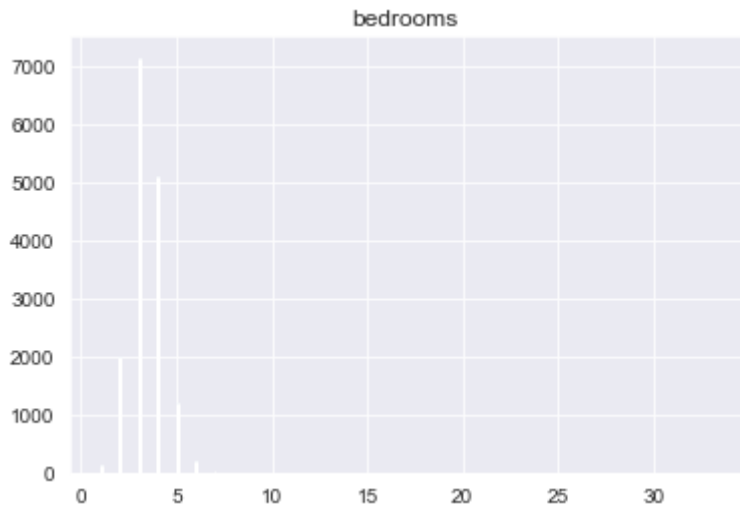
Normalization/Log Transformations



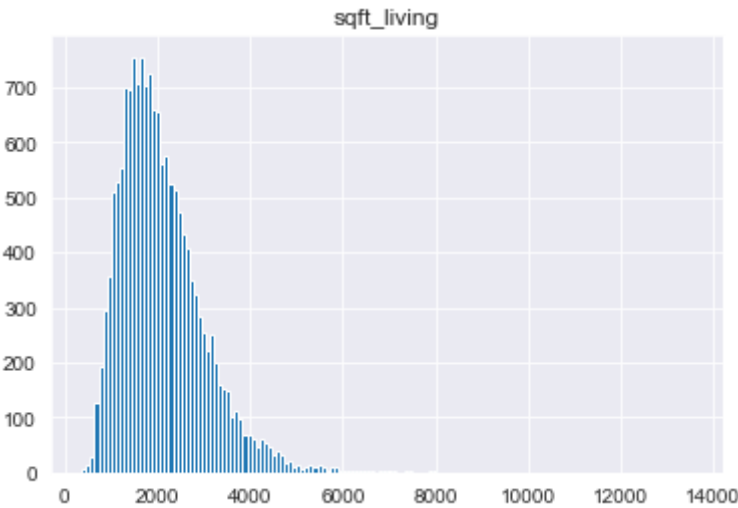
Using histograms, continuous variables are checked for normality and outliers

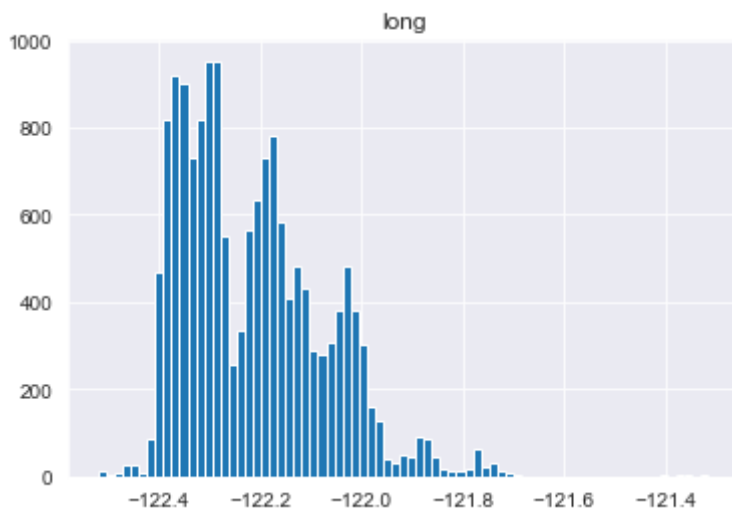
```

In [19]: continuous = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'yr_built', 'lat',
for var in continuous:
    df[var].hist(bins='auto')
    plt.title(var)
    plt.show()

```








```
In [20]: # Bedroom outliers
df.drop(df.loc[df['bedrooms'] > 8].index, inplace=True)
df['bedrooms'].hist(bins = 'auto')
```

Out[20]: <AxesSubplot:>



```
In [21]: # Log Transformations
df['log_long'] = df['long'].map(lambda x: np.log(x*-1))
df.drop('long', axis = 1, inplace = True)
df['log_lat'] = df['lat'].map(lambda x: np.log(x))
df.drop('lat', axis = 1, inplace = True)
df['log_yr_built'] = df['yr_built'].map(lambda x: np.log(x))
df.drop('yr_built', axis = 1, inplace = True)
df['log_sqft_living'] = df['sqft_living'].map(lambda x: np.log(x))
df.drop('sqft_living', axis = 1, inplace = True)
df['log_bed'] = df['bedrooms'].map(lambda x: np.log(x))
df.drop('bedrooms', axis = 1, inplace = True)
df['log_price'] = df['price'].map(lambda x: np.log(x))
df.drop('price', axis = 1, inplace = True)
```

```
In [22]: df.head()
```

```
Out[22]:
```

	bathrooms	waterfront	condition_2	condition_3	condition_4	condition_5	grade_4	grade_5	g
1	2.25	0.0	0	1	0	0	0	0	
3	3.00	0.0	0	0	0	1	0	0	

	bathrooms	waterfront	condition_2	condition_3	condition_4	condition_5	grade_4	grade_5	g
4	2.00	0.0	0	1	0	0	0	0	
5	4.50	0.0	0	1	0	0	0	0	
6	2.25	0.0	0	1	0	0	0	0	

5 rows × 92 columns

```
In [23]: lin_reg_4 = LinearRegression()
y = df['log_price'].map(lambda x: np.log(x))
X = df.drop(['log_price'], axis = 1)
lin_reg_4.fit(X, y)
print('R-Squared: ', lin_reg_4.score(X, y))
```

R-Squared: 0.8666057935678042

With a model that represents 86% of the variation of data, We can move forward to model evaluation.

Model Evaluation

Using Scikitlearn to randomly split the data into a train group and a test group, the model can be evaluated for accuracy

```
In [24]: from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as mse
```

```
In [25]: X = df.drop('log_price', axis = 1)
y = df['log_price']
```

```
In [26]: X_train, X_test, y_train, y_test = train_test_split(X,y)
```

```
In [27]: linreg = LinearRegression()
linreg.fit(X_train, y_train)
y_hat_train = linreg.predict(X_train)
y_hat_test = linreg.predict(X_test)
print('R-Squared: ', linreg.score(X_train, y_train))
```

R-Squared: 0.8694121573436995

```
In [28]: train_mse = mse(y_train, y_hat_train)
test_mse = mse(y_test, y_hat_test)
print('Train MSE: ', train_mse)
print('Test MSE: ', test_mse)
```

Train MSE: 0.03597150394098349

Test MSE: 0.03731187808602213


The mean squared error of the test data is actually lower than the mean squared error of the training data. This is a good sign that our model is properly fitted.

Lastly, the distribution of residuals will be analyzed.


```
In [29]: residuals = y_test - linreg.predict(X_test)
```

```
In [30]: sns.histplot(residuals);
```

```
plt.title('Residual Histogram');
```



```
In [31]: import scipy.stats as stats
model = sm.OLS(y_test, sm.add_constant(X_test)).fit()
sm_residuals = model.resid
fig = sm.graphics.qqplot(sm_residuals, dist=stats.norm, line='45', fit=True)
```



The residuals of the model are not perfectly normal. this may cause some issues with price predictions at the upper and lower levels of our distribution.

Deployment

Using the pickle module, we can store our model and use it in a leaner notebook or .py file to predict the value of a house in King County.

```
In [32]: import pickle
with open('deployment/regression_model.pickle', 'wb') as f:
    pickle.dump(linreg, f)
```

```
In [33]: house_df = pd.DataFrame(np.zeros((1, len(X.columns))), columns = X.columns)
with open('deployment/house_data.pickle', 'wb') as f:
    pickle.dump(house_df, f)
```

Using our regression model and a datadrame template to be filled out and fed to the model, I

created `house_price_prediction.py` which can be run from the terminal to predict real estate prices in King County

Conclusions

After cleaning and normalizing data on real estate prices in King County, A multiple linear regression model was created that captures 86% of our data's variance.

The dataset is suitably linear, and its residuals follow a near-normal distribution, but there is some heteroscedasticity in the data that suggests some inaccuracy.

This model was exported to `house_price_prediction.py` so it can be used in the backend of a prediction software for a real estate company to automate the initial appraisal of houses in King County.

Future Work

- Create GUI that takes data and delivers a price estimate.
- Explore using multiple models for different locations to achieve higher accuracy
- Modify function to accept an address as an input rather than taking zipcode, latitude, and longitude separately